

Drink Dataset

Data Resource : <https://github.com/fivethirtyeight/data/blob/master/alcohol-consumption/drinks.csv>
(<https://github.com/fivethirtyeight/data/blob/master/alcohol-consumption/drinks.csv>)

Story Behind : <https://fivethirtyeight.com/features/dear-mona-followup-where-do-people-drink-the-most-beer-wine-and-spirits/> (<https://fivethirtyeight.com/features/dear-mona-followup-where-do-people-drink-the-most-beer-wine-and-spirits/>)

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import Ridge
from sklearn.preprocessing import PolynomialFeatures
```

```
In [3]: df= pd.read_csv('drinks.csv')
```

```
In [4]: df.head()
```

Out[4]:

	country	beer_servings	spirit_servings	wine_servings	total_litres_of_pure_alcohol	continent
0	Afghanistan	0	0	0	0.0	Asia
1	Albania	89	132	54	4.9	Europe
2	Algeria	25	0	14	0.7	Africa
3	Andorra	245	138	312	12.4	Europe
4	Angola	217	57	45	5.9	Africa

```
In [5]: df.dtypes
```

```
Out[5]: country          object
beer_servings          int64
spirit_servings        int64
wine_servings          int64
total_litres_of_pure_alcohol  float64
continent              object
dtype: object
```

Exploratory Data Analysis

```
In [6]: df[['wine_servings', 'continent']].groupby('continent').sum()
```

Out[6]:

wine_servings	
continent	
Africa	862
Asia	399
Europe	6400
North America	564
Oceania	570
South America	749

```
In [7]: df[["beer_servings", "continent"]].groupby(["continent"]).describe()
```

Out[7]:

beer_servings									
	count	mean	std	min	25%	50%	75%	max	
continent									
Africa	53.0	61.471698	80.557816	0.0	15.00	32.0	76.00	376.0	
Asia	44.0	37.045455	49.469725	0.0	4.25	17.5	60.50	247.0	
Europe	45.0	193.777778	99.631569	0.0	127.00	219.0	270.00	361.0	
North America	23.0	145.434783	79.621163	1.0	80.00	143.0	198.00	285.0	
Oceania	16.0	89.687500	96.641412	0.0	21.00	52.5	125.75	306.0	
South America	12.0	175.083333	65.242845	93.0	129.50	162.5	198.00	333.0	


```
In [8]: fig, ax = plt.subplots(nrows=4, ncols=1, figsize=(10,24))


column = ['beer_servings', 'spirit_servings', 'wine_servings', 'total_litre
s_of_pure_alcohol']
category = [df[i].value_counts() for i in column]
r = 0
for j in range(4):

    axes = ax[j]

    sns.boxplot(x=column[r], y="continent", data=df, ax=axes)
    axes.set_xlabel('')
    axes.set_title(column[r])
    r+=1

plt.show()
```

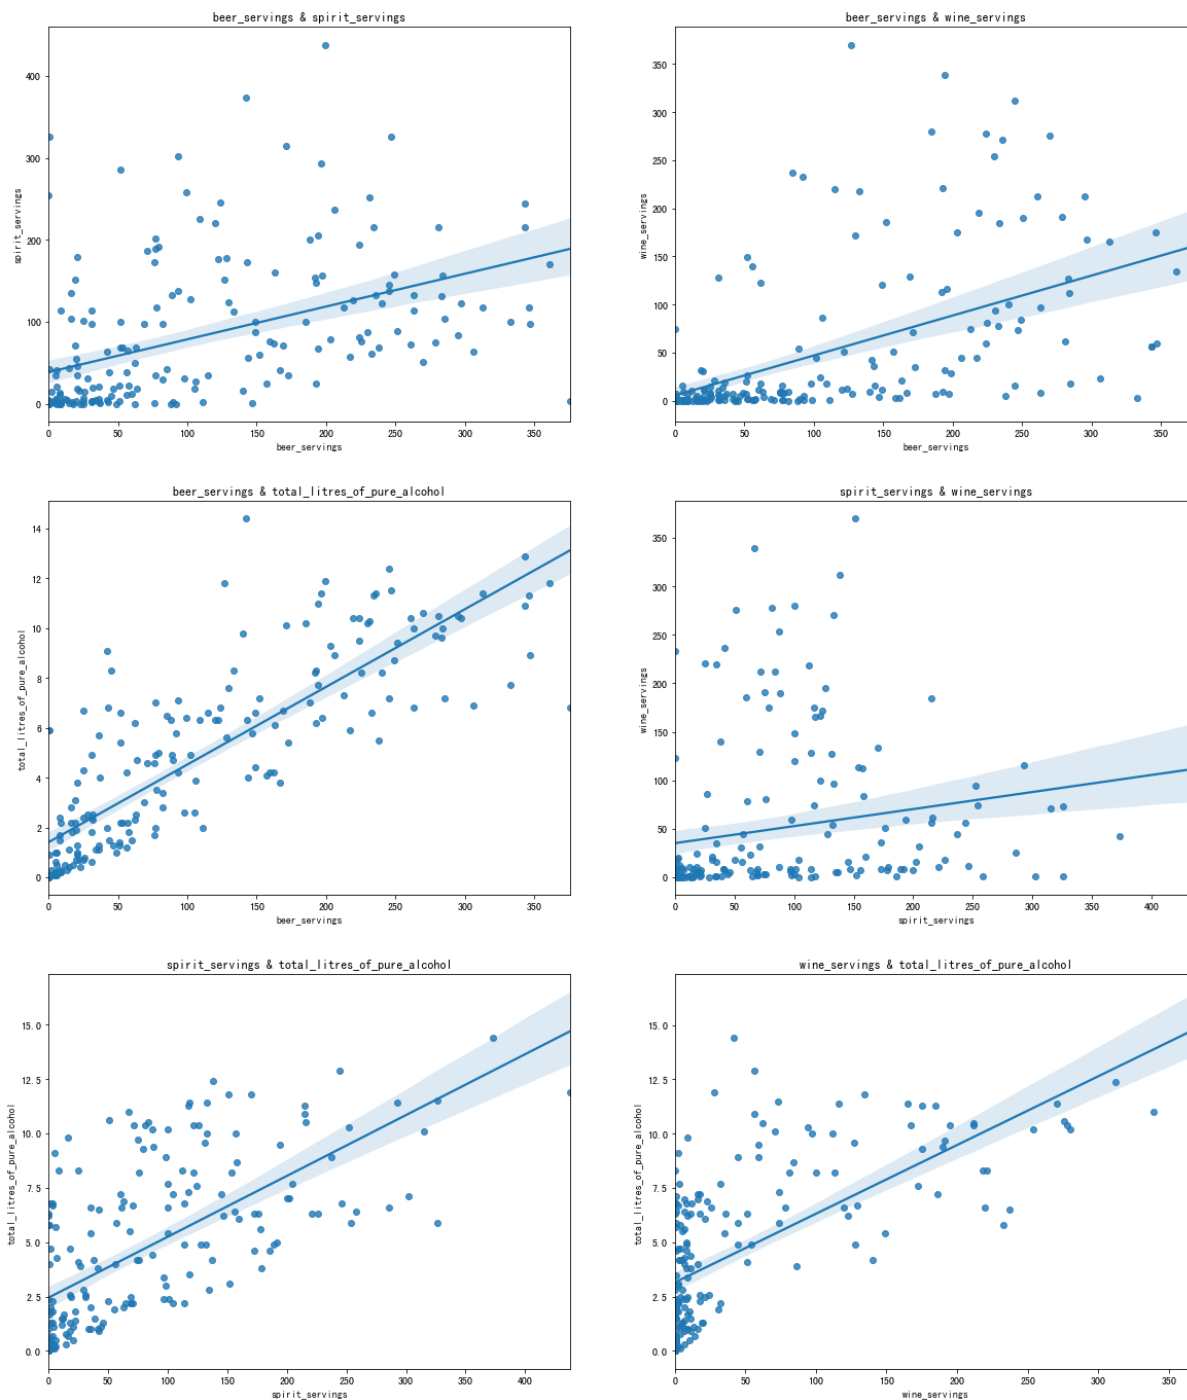




```
In [9]: fig, ax = plt.subplots(nrows=3, ncols=2, figsize=(20,24))
column = ['beer_servings', 'spirit_servings', 'wine_servings', 'total_litre
s_of_pure_alcohol']
category = [df[i].value_counts() for i in column]
axe = [ax[i][j] for i in range(3) for j in range(2)]
r = 0

for i in range(3):
    for j in range(i+1,4):
        axes = axe[r]
        sns.regplot(x=column[i],y=column[j],data=df,ax=axes)
        axes.set_title(f'{column[i]} & {column[j]}')
        r+=1

plt.show()
```



Modeling

```
In [10]: print('LinearRegression - Wine and Total')
print()
lm = LinearRegression()
x=df[['wine_servings']]
y=df[['total_litres_of_pure_alcohol']]
lm.fit(x,y)
lm.score(x,y)
print(f'Total = ({str(lm.coef_[0][0])[:6]} * Wine) + {str(lm.intercept_
[0])[:6]}')
print('R^2 = ',lm.score(x,y))
```

LinearRegression - Wine and Total

Total = (0.0316 * Wine) + 3.1540
R^2 = 0.4456875459787605

```
In [11]: print('LinearRegression - Beer and Total')
print()
lm = LinearRegression()
x=df[['beer_servings']]
y=df[['total_litres_of_pure_alcohol']]
lm.fit(x,y)
lm.score(x,y)
print(f'Total = ({str(lm.coef_[0][0])[:6]} * Beer) + {str(lm.intercept_
[0])[:6]}')
print('R^2 = ',lm.score(x,y))
```

LinearRegression - Beer and Total

Total = (0.0311 * Beer) + 1.4067
R^2 = 0.6986262119487155

```
In [12]: print('LinearRegression - Spirit and Total')
print()
lm = LinearRegression()
x=df[['spirit_servings']]
y=df[['total_litres_of_pure_alcohol']]
lm.fit(x,y)
lm.score(x,y)
print(f'Total = ({str(lm.coef_[0][0])[:6]} * Spirit) + {str(lm.intercept_
[0])[:6]}')
print('R^2 = ',lm.score(x,y))
```

LinearRegression - Spirit and Total

Total = (0.0279 * Spirit) + 2.4497
R^2 = 0.4289833103726812


```
In [13]: print('LinearRegression - Wine + Beer + Spirit and Total')
print()
Z=df[['wine_servings',"beer_servings", "spirit_servings" ]]
y=df[['total_litres_of_pure_alcohol']]


Z_train , Z_test , y_train, y_test = train_test_split(Z , y , test_size=
0.3, random_state=0)
lm1 = LinearRegression()
lm1.fit(Z_train , y_train)
print(f'Total = ({str(lm1.coef_[0][0])[:6]} * Wine) + ({str(lm1.coef_[0]
[1])[:6]} * Beer) + ({str(lm1.coef_[0][2])[:6]} * Spirit) + {str(lm1.int
ercept_[0])[:6]}')

print("R^2 = ",lm1.score(Z_test , y_test))

lm1.predict(Z_test)
plt.scatter(y_test,lm1.predict(Z_test))
plt.title('predict and actual')
plt.show()
```

LinearRegression - Wine + Beer + Spirit and Total

Total = (0.0165 * Wine) + (0.0181 * Beer) + (0.0169 * Spirit) + 0.4779
R^2 = 0.7549608066333018




```
In [14]: print('PipeLine - Standard , Polynomial(degree=3) , LinearRegression')
print()
Input = [('scale',StandardScaler()), ('polynomial', PolynomialFeatures(degree=3, include_bias=False)), ('model',LinearRegression())]
pipe = Pipeline(Input)

pipe.fit(Z_train,y_train)
print("R^2 = ",pipe.score(Z_test,y_test))
plt.scatter(pipe.predict(Z_test),y_test)
plt.title('predict and actual')
plt.show()
```

PipeLine - Standard , Polynomial(degree=3) , LinearRegression

R^2 = 0.7746820927843575




```
In [15]: print('RidgeModel - alpha(0.1)')
print()
RidgeModel=Ridge(alpha=0.1)
RidgeModel.fit(Z_train,y_train)
print("R^2 = ",RidgeModel.score(Z_test,y_test))
print(f'Total = ({str(RidgeModel.coef_[0][0])[:6]} * Wine) + ({str(RidgeModel.coef_[0][1])[:6]} * Beer) + ({str(RidgeModel.coef_[0][2])[:6]} * Spirit) + {str(RidgeModel.intercept_[0])[:6]}')
plt.scatter(RidgeModel.predict(Z_test),y_test)
plt.title('predict and actual')
plt.show()
```

#Ridge(alpha=X) alpha越大懲罰越大

RidgeModel - alpha(0.1)

R^2 = 0.7549608184502712

Total = (0.0165 * Wine) + (0.0181 * Beer) + (0.0169 * Spirit) + 0.4779



```
In [16]: print('Polynomial(degree=3) + RidgeModel - alpha(0.1) ')
print()
pr=PolynomialFeatures(degree=3)
Z_train_pr=pr.fit_transform(Z_train)
Z_test_pr=pr.fit_transform(Z_test)

RidgeModel=Ridge(alpha=0.2)
RidgeModel.fit(Z_train_pr,y_train)
print("R^2 = ",RidgeModel.score(Z_test_pr,y_test))

plt.scatter(RidgeModel.predict(Z_test_pr),y_test)
plt.title('predict and actual')
plt.show()
```

Polynomial(degree=3) + RidgeModel - alpha(0.1)

R^2 = 0.7746824248087052

/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-packages/sklearn/linear_model/_ridge.py:147: LinAlgWarning: Ill-conditioned matrix (rcond=6.84e-18): result may not be accurate.
return linalg.solve(A, Xy, sym_pos=True,

