

Heart Attack Analysis & Prediction Dataset

Data Resource : <https://www.kaggle.com/rashikrahmanpritom/heart-attack-analysis-prediction-dataset>
(<https://www.kaggle.com/rashikrahmanpritom/heart-attack-analysis-prediction-dataset>)

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Import Data

```
In [2]: Heart = pd.read_csv('heart.csv')
```

```
In [3]: Heart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null   int64
 1   sex         303 non-null   int64
 2   cp          303 non-null   int64
 3   trtbps      303 non-null   int64
 4   chol        303 non-null   int64
 5   fbs         303 non-null   int64
 6   restecg     303 non-null   int64
 7   thalachh    303 non-null   int64
 8   exng        303 non-null   int64
 9   oldpeak     303 non-null   float64
10   slp         303 non-null   int64
11   caa         303 non-null   int64
12   thall       303 non-null   int64
13   output      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
In [4]: Heart.describe()
```

```
Out[4]:
```

	age	sex	cp	trtbps	chol	fbs	restecg	th
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.

```
In [5]: Heart.nunique()
```

```
Out[5]: age          41
sex            2
cp             4
trtbps        49
chol         152
fbs           2
restecg       3
thalachh     91
exng          2
oldpeak      40
slp           3
caa           5
thall        4
output        2
dtype: int64
```

Exploratory Data Analysis

```
In [6]: Heart.head(10)
```

```
Out[6]:
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
5	57	1	0	140	192	0	1	148	0	0.4	1	0	1	1
6	56	0	1	140	294	0	0	153	0	1.3	1	0	2	1
7	44	1	1	120	263	0	1	173	0	0.0	2	0	3	1
8	52	1	2	172	199	1	1	162	0	0.5	2	0	3	1
9	57	1	2	150	168	0	1	174	0	1.6	2	0	2	1

```
In [7]: cat = ['sex', 'cp', 'fbs', 'restecg', 'exng', 'slp', 'caa', 'thall']
len(cat)
```

```
Out[7]: 8
```

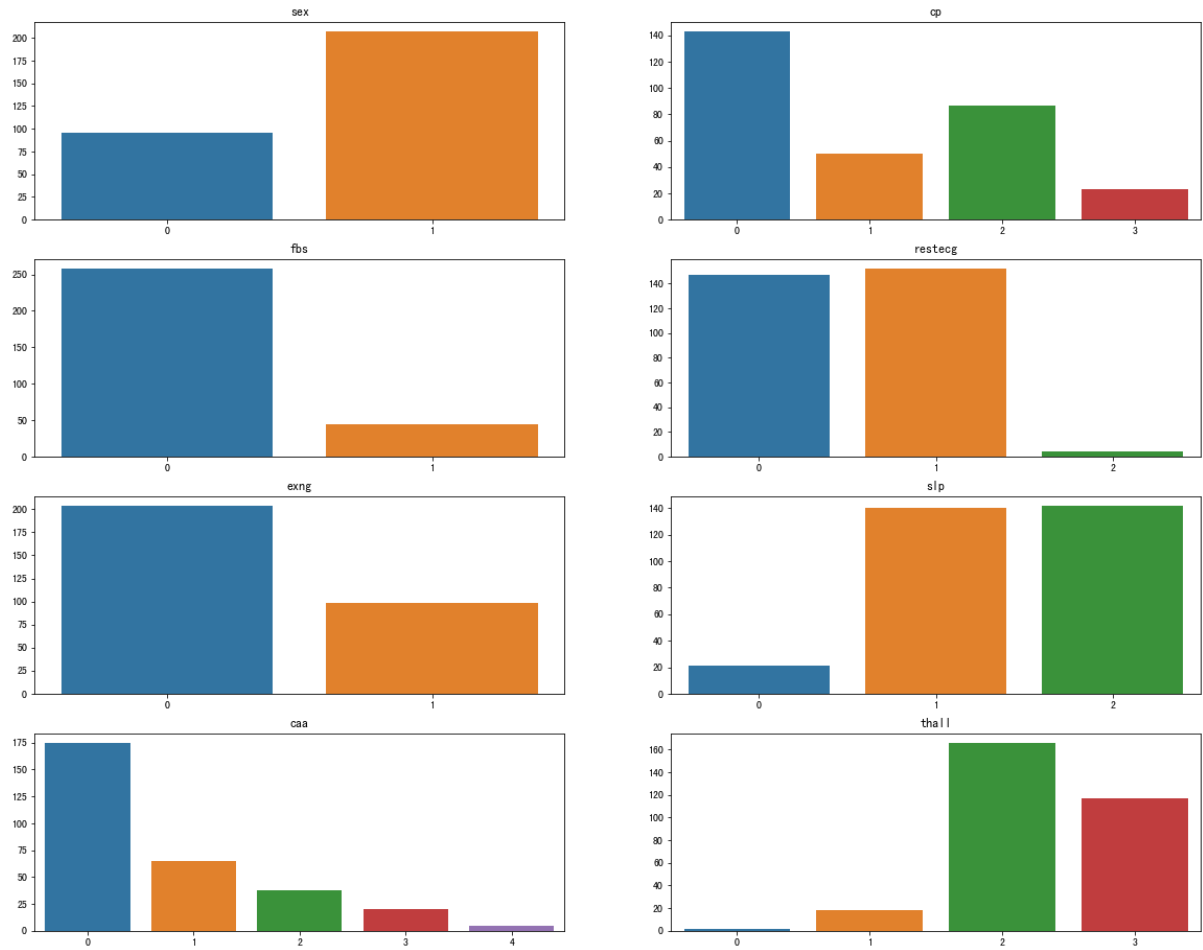
```
In [8]: Heart['sex'].value_counts().index
```

```
Out[8]: Int64Index([1, 0], dtype='int64')
```

```

In [9]: fig, ax = plt.subplots(nrows=4, ncols=2, figsize=(20,16))
r = 0
for i in range(4):
    for j in range(2):
        category = [Heart[i].value_counts() for i in cat]
        axes = ax[i][j]
        sns.barplot(x=category[r].index, y=category[r].values, ax=axes)
        axes.set_title(cat[r])
        r+=1
plt.show()

```



About this dataset

- sex : sex of the patient (0 = Female; 1 = Male)
- exang: exercise induced angina (1 = yes; 0 = no)
- ca: number of major vessels (0-3)
- cp : Chest Pain type chest pain type
 - Value 1: typical angina
 - Value 2: atypical angina
 - Value 3: non-anginal pain
 - Value 4: asymptomatic
- fbs : (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)
- rest_ecg : resting electrocardiographic results
 - Value 0: normal
 - Value 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV)
 - Value 2: showing probable or definite left ventricular hypertrophy by Estes' criteria


Analysis

- ex : 68% Man ($207 / (207 + 96)$)
- cp : 47% of typical angina ($143 / (143 + 87 + 50 + 23)$)
- fbs : 85% of people have fasting blood sugar \leq 120 mg/dl ($45 / (258 + 45)$)
- rest_ecg : normal & ST-T wave abnormality are 98% ($(152 + 147) / (152 + 147 + 4)$)
- exang: 67% people exercise "without" induced angina ($204 / (204 + 99)$)

As we know 55% of the people have high risk of Heart Attack

```
In [10]: target = Heart.output.value_counts()  
sns.barplot(x=target.index,y=target.values)
```

```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcc97002dc0>
```



Correlation

```
In [11]: Output = pd.DataFrame(Heart.corr()['output'].sort_values(ascending=False  
)  
Heart.corr()['output'].sort_values(ascending=False))
```

```
Out[11]: output      1.000000  
cp      0.433798  
thalachh 0.421741  
slp      0.345877  
restecg  0.137230  
fbs     -0.028046  
chol    -0.085239  
trtbps  -0.144931  
age     -0.225439  
sex     -0.280937  
thall   -0.344029  
caa     -0.391724  
oldpeak -0.430696  
exng    -0.436757  
Name: output, dtype: float64
```

```
In [12]: sns.heatmap(Output)
```


```
/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-packages/matplotlib/backends/backend_agg.py:214: RuntimeWarning: Glyph 8722 missing from current font.
```

```
font.set_text(s, 0.0, flags=flags)
```

```
/Library/Frameworks/Python.framework/Versions/3.8/lib/python3.8/site-packages/matplotlib/backends/backend_agg.py:183: RuntimeWarning: Glyph 8722 missing from current font.
```

```
font.set_text(s, 0, flags=flags)
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7fcc97eeeb80>
```



Data Preprocessing

```
In [13]: from sklearn.model_selection import train_test_split
```

```
In [14]: train = Heart.iloc[:, :-1]
test = Heart.iloc[:, -1:]
```

```
In [15]: X_train, X_test, y_train, y_test = train_test_split(train, test, test_size=0.3, random_state=3)
```

In [16]: X_train

Out[16]:

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall
197	67	1	0	125	254	1	1	163	0	0.2	1	2	3
15	50	0	2	120	219	0	1	158	0	1.6	1	0	2
227	35	1	0	120	198	0	1	130	1	1.6	1	0	3
146	44	0	2	118	242	0	1	149	0	0.3	1	1	2
115	37	0	2	120	215	0	1	170	0	0.0	2	0	2
...
277	57	1	1	124	261	0	1	141	0	0.3	2	0	3
256	58	1	0	128	259	0	0	130	1	3.0	1	2	3
131	49	0	1	134	271	0	1	162	0	0.0	1	0	2
249	69	1	2	140	254	0	0	146	0	2.0	1	3	3
152	64	1	3	170	227	0	0	155	0	0.6	1	0	3

212 rows × 13 columns

In [17]: `print(X_train.shape, X_test.shape)`
`print(y_train.shape, y_test.shape)`

(212, 13) (91, 13)
 (212, 1) (91, 1)

In [18]: `from sklearn.preprocessing import StandardScaler`

In [19]: `scaler = StandardScaler()`
`X_train_raw = scaler.fit_transform(X_train)`
`X_test_raw = scaler.transform(X_test)`

In [20]: `X_train = pd.DataFrame(X_train_raw, columns=X_train.columns, index=X_train.index)`
`X_test = pd.DataFrame(X_test_raw, columns=X_test.columns, index=X_test.index)`


```
In [21]: X_train
```

```
Out[21]:
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	e
197	1.397056	0.650245	-0.901439	-0.356084	0.154583	2.617604	0.878509	0.652301	-0.717
15	-0.480271	-1.537881	1.068716	-0.656362	-0.492710	-0.382029	0.878509	0.429184	-0.717
227	-2.136735	0.650245	-0.901439	-0.656362	-0.881086	-0.382029	0.878509	-0.820270	1.394
146	-1.142857	-1.537881	1.068716	-0.776473	-0.067346	-0.382029	0.878509	0.027574	-0.717
115	-1.915873	-1.537881	1.068716	-0.656362	-0.566687	-0.382029	0.878509	0.964665	-0.717
...
277	0.292746	0.650245	0.083639	-0.416140	0.284041	-0.382029	0.878509	-0.329413	-0.717
256	0.403177	0.650245	-0.901439	-0.175918	0.247053	-0.382029	-1.021939	-0.820270	1.394
131	-0.590702	-1.537881	0.083639	0.184416	0.468982	-0.382029	0.878509	0.607678	-0.717
249	1.617918	0.650245	1.068716	0.544749	0.154583	-0.382029	-1.021939	-0.106296	-0.717
152	1.065763	0.650245	2.053793	2.346417	-0.344758	-0.382029	-1.021939	0.295314	-0.717

212 rows × 13 columns

Modeling Using LogisticRegression

```
In [22]: from sklearn.linear_model import LogisticRegression
```

```
In [23]: lr = LogisticRegression()
```

```
In [24]: lr.fit(X_train,y_train.values.ravel())
y_pred = lr.predict(X_test)
```

```
In [25]: (y_pred == y_test.values.ravel()).sum() / len(y_pred)
```

```
Out[25]: 0.8791208791208791
```

```
In [26]: from sklearn.metrics import accuracy_score
```

```
In [27]: accuracy_score(y_pred, y_test)
```

```
Out[27]: 0.8791208791208791
```